

# ML + LLM Pipeline Orchestration Documentation

## Overview:

This documentation outlines a complete orchestration project integrating a Machine Learning (ML) pipeline with a Retrieval-Augmented Generation (RAG) system. The project uses modern tools like Apache Airflow, MLflow, FastAPI, and Docker for seamless automation, model serving, and document-based LLM querying.

## Part A: ML Pipeline (Wine Quality Prediction)

**Objective:** Automate the training, evaluation, logging, and deployment of a classification model for predicting wine quality.

### Tech Stack:

- **Airflow:** Workflow orchestration
- **Pandas:** Data preprocessing and feature engineering
- **Scikit-learn:** ML model development
- **MLflow:** Model tracking and experiment logging
- **FastAPI:** REST API for predictions
- **Docker:** Containerization and orchestration

### Directory Structure:

```
ml_pipeline/
├── dags/
│   └── ml_pipeline_dag.py
├── scripts/
│   ├── ingest_data.py
│   ├── preprocess.py
│   ├── feature_engineering.py
│   ├── train_model.py
│   ├── evaluate_model.py
│   └── log_to_mlflow.py
├── model_service/
│   ├── Dockerfile
│   ├── main.py
│   └── requirements.txt
├── data/
└── requirements.txt
```

### ML Workflow:

1. **Ingest:** Load dataset (winequality-red.csv)
2. **Preprocess:** Clean, normalize and handle missing values
3. **Feature Engineering:** Generate additional features
4. **Train:** Train a classification model (e.g., RandomForest)
5. **Evaluate:** Calculate metrics (accuracy, precision, etc.)
6. **Log to MLflow:** Store model and metadata
7. **Serve:** Expose prediction API via FastAPI

### FastAPI ML Inference Test:

```
curl -X POST http://localhost:8000/predict \
-H "Content-Type: application/json" \
```

```
-d '{"data": [[7.4, 0.7, 0, 1.9, 0.076, 11, 34, 0.9978, 3.51, 0.56, 9.4]]}'
```

**Response:**

```
{"predictions": [5]}
```

## Part B: RAG LLM Pipeline (Mini POC)

**Objective:** Retrieve context from PDF documents and answer user queries using embedded vectors and a lightweight transformer model.

**Tech Stack:**

- **FAISS:** Efficient similarity search over embeddings
- **DistilBERT:** For text embeddings
- **FastAPI:** For serving the RAG system
- **Docker:** For containerization

**Directory Structure:**

```
rag_pipeline/
├── app/
│   ├── main.py
│   ├── vector_store.py
│   ├── requirements.txt
│   └── Dockerfile
├── docs/
├── embeddings/
└── requirements.txt
```

**Workflow:**

1. Upload PDF to docs/
2. Split and embed content
3. Store in FAISS index
4. Query via REST API
5. Return LLM-generated answer using context

**FastAPI RAG Query Test:**

```
curl -X GET http://localhost:8001/query?q=What is the project about?
```

**Response:**

```
{
  "query": "What is the project about?",
  "answer": "This project demonstrates a full-cycle ML pipeline and a simple LLM-based RAG system..."
}
```

## Project Setup Instructions

**Step-by-Step Setup:****# Clone Repository**

```
$ git clone <your-repo-url>
```

```
$ cd ml_llm_pipeline_template
```

**# Optional Cleanup**

```
$ docker system prune -af --volumes
```

## # Build Images

```
$ docker-compose build --no-cache
```

## # Start All Services

```
$ docker-compose up -d
```

## Docker Services and Ports:

Service	Purpose	Port
airflow	DAG execution	<a href="http://localhost:8080">http://localhost:8080</a>
mlflow	Model tracking	<a href="http://localhost:5000">http://localhost:5000</a>
model_service	ML API	<a href="http://localhost:8000/docs">http://localhost:8000/docs</a>
rag_service	RAG API	<a href="http://localhost:8001/docs">http://localhost:8001/docs</a>
spark	(Optional) Spark support	—
postgres	Airflow backend	—

## Retrain & Register Model:

```
$ docker-compose exec airflow python /scripts/train_model.py
```

```
$ docker-compose restart model_service
```

## Airflow DAG

- **DAG ID:** wine\_quality\_pipeline
- **Schedule:** Weekly (@weekly)
- Access via: <http://localhost:8080>

## MLflow Tracking

- **Access URL:** <http://localhost:5000>
- Tracks: Metrics, Parameters, Registered Models, Artifacts

## Challenges Faced

- Airflow import/path configurations
- MLflow logging context inside Airflow
- PySpark compatibility in Docker
- Chunking and FAISS index freshness
- Docker volume and port management

## Improvements & Scope

- Add PySpark back for scalable preprocessing
- CI/CD with GitHub Actions
- Real-time alerts via email/Slack
- Use Mistral-7B or TinyLLaMA for better RAG response
- Interactive UI for LLM queries

## Conclusion

This orchestration project successfully demonstrates how to:

- Automate an ML lifecycle
- Log and manage models with MLflow
- Deploy scalable APIs with FastAPI
- Implement a functional RAG pipeline

With modularity and Dockerization, the system is easily extensible to production-grade workflows.